

ECVOL-bench: An Open, Identity-Controlled Benchmark for Multimodal Earnings-Call Volatility Prediction

(anonymized working draft for comments — please do not cite or circulate)

Anonymous Author(s)

Draft compiled August 26, 2026

Abstract

A decade of work predicts post-earnings-call stock volatility from the text and audio of quarterly conference calls. We argue that the evaluation substrate this literature shares cannot support its claims: the canonical benchmarks are scored with a zero-day embargo against 30-day targets, with 44–80% of test calls from companies also present in training, on label files inherited unchanged from paper to paper, and under a horizon convention that downstream work misattributes *pending: regenerated by ecvol audit substrate*, Table 1]. We release ECVOL-bench, an open multimodal benchmark over FinCall-Surprise and MAEC that repairs this substrate: reconstructed company identity and call dates for a corpus that ships without them (92.9% of calls; 95.3% of earnings calls; 60/60 manual audit), three target families (level log-realised-volatility, change Δv , HAR-residual) under *both* the calendar-day and trading-day horizon conventions, temporal and ticker-disjoint splits with a 30-trading-day embargo, SHA-256 manifests, strong econometric baselines, and an identity-control suite (linear ticker probes, same-ticker and global content shuffles). Run over modern open-weight text and audio representations, the benchmark returns a consistent answer: text embeddings encode *which company* is speaking (89.5% ticker-probe accuracy; within-ticker transcript shuffles leave performance unchanged), frozen audio content is inert beyond past volatility (its apparent gains survive even a global audio shuffle), and fusion never significantly beats the best unimodal model—under honest controls, apparent content signal in this task family is overwhelmingly company identity plus past volatility. *pending: We then re-run the field’s five publicly runnable models on their own data under these controls (Table 7), re-evaluate the audio ladder on clean 2025 recordings, and evaluate the frozen pipeline on post-LLM-cutoff calls.* All code, configs, manifests, reconstructed metadata, split files, targets, derived features, and result tables are released.

1 Introduction

Earnings conference calls are one of the few recurring events at which management speaks at length, semi-

spontaneously, about a public company—in both *what* is said (the transcript) and *how* it is said (the audio). A line of research beginning with Qin and Yang (2019) and continuing through HTML (Yang et al., 2020), MAEC (Li et al., 2020), KeFVP (Niu et al., 2023), and ECC Analyzer (Cao et al., 2024a) claims that multimodal models over this data predict post-call stock volatility better than classical methods.

Two recent observations cast doubt on how much of that progress is real. First, Yu et al. (2024) (“Same Company, Same Signal”) show that transcript representations used across this literature predominantly encode *company identity* rather than content: same-ticker transcripts cluster tightly, and training-free baselines built on a company’s own past volatility beat all transcript-based models on their evaluation. Because volatility is strongly autocorrelated and company-specific, a model that learns “which company is this” inherits most of the apparent skill. Second, a 2025 survey of financial NLP (Tatarinov et al., 2025) finds that only $\sim 14\%$ of reported results reproduce exactly, attributing the gap to dead data links, API drift, and contamination.

This paper’s premise is that both problems are properties of the *evaluation substrate* the field shares, and that the substrate can be repaired. The lineage above evaluates on a small number of label and split files that are passed from paper to paper essentially unchanged. Those files (§3.1) impose no embargo between training and test windows even though targets run 30 days past each call; they place most test calls’ companies in the training set, and no paper reports the ticker-disjoint condition; they carry degenerate label cells; and they encode a horizon convention (calendar days) that the downstream literature describes as trading days. On such a substrate a model cannot be told apart from a memory of which company it is looking at. Meanwhile the corpus with the best license and the only raw audio, FinCall-Surprise (Shu et al., 2025), ships with *no* ticker, company, date, or speaker metadata at all—so the identity controls that would expose the confound cannot even be run on it without first reconstructing the identity the dataset omits.

We therefore build and release ECVOL-bench, an open benchmark that supplies what the substrate lacks, and we use it to re-examine the task’s central claims. The

benchmark is the contribution; the empirical results are the evidence for why it is needed. Concretely, we ask four research questions:

RQ1 Does earnings-call *content* (text and/or audio) add predictive signal for post-call volatility beyond past-volatility persistence (HAR-RV, GARCH)?

RQ2 Is multimodal (audio + text) better than the best unimodal model, after controlling for past volatility?

RQ3 Do the published models of this lineage, re-run on their own data, retain their reported advantage under embargoed, ticker-disjoint evaluation?

RQ4 Do any positive results survive (a) ticker-identity controls, (b) ticker-disjoint splits, and (c) evaluation on post-LLM-knowledge-cutoff calls?

The study was pre-registered so that no single outcome strands the work; the modelling stages, controls, and confirmatory tests were fixed before any model result existed and are recorded in a public decision log. RQ1, RQ2, and RQ4(a–b) are answered here; RQ3 and RQ4(c) are the subject of the reproduction and post-cutoff studies that this draft still owes (marked *[pending: like this]*).

Contributions.

1. **ECVOL-bench**, an open multimodal volatility benchmark over FinCall-Surprise (primary) and MAEC (secondary): reconstructed identity and dates for FinCall-Surprise (92.9% of calls; 95.3% of earnings calls; 60/60 manual audit), three target families (level v , change Δv , HAR-residual) under both the calendar-day and trading-day horizon conventions, temporal and ticker-disjoint splits with a 30-trading-day embargo, SHA-256 manifests, and a reproducible harness—released in full (§3).
2. **A substrate audit** of the label and split files the prior literature shares: embargo, ticker overlap, label defects, and the horizon convention, regenerable from one command (§3.1, Table 1).
3. **Honest baselines and an identity-control suite**—persistence, EWMA, HAR-RV, GARCH(1,1), and ticker-fixed-effect models for every horizon, split, and target; identity probes; transcript- and audio-shuffle controls; ticker-disjoint evaluation; a gender-confound analysis—and what they show when modern open-weight text and audio representations are run through them on consumer hardware (§6).
4. **A reproduction study** of the five publicly runnable models in this lineage, each re-run on its own original data under the new controls, together with a code-availability audit of the fifteen that cannot be run *[pending: Phase 6R; Table 7]*.

5. **Post-cutoff evaluation**: the audio ladder re-run on clean 2025 recordings, and the frozen pipeline evaluated on calls held after the knowledge cutoff of every model used *[pending: Phase 7 and T9.4]*.

2 Related Work

The benchmark lineage. Qin and Yang (2019) introduced text+audio volatility prediction and the 572-call EarningsCall dataset, establishing the 3/7/15/30-day log-realised-volatility target family that the field still uses (defined over *calendar* days in the original, a detail that matters in §3.1). MAEC (Li et al., 2020) scaled this to ~3,443 sentence-aligned multimodal calls over the S&P 1500. HTML (Yang et al., 2020) introduced a hierarchical transformer with multi-task learning across horizons; VolTAGE (Sawhney et al., 2020) added graph convolution over inter-stock relations with text/audio fusion; DialogueGAT (Sang and Bao, 2022) modelled speaker/dialogue structure with graph attention. Ke-FVP (Niu et al., 2023) combined knowledge-enhanced pre-training with a price-conditioned time-series module and is the standard SOTA reference. More recent work moves toward LLMs: ECC Analyzer (Cao et al., 2024a) uses hierarchical LLM summarisation with a RAG “question bank” and wav2vec2 features; RiskLabs (Cao et al., 2024b) fuses calls with news and market series; AT-FinGPT (Liu et al., 2025) is an audio-text LLM. Numeral-aware variants (NumHTML (Yang et al., 2022), NAM/ECNum (Chen et al., 2021a), GNAVol (Shi et al., 2023)) and Q&A-structure models (Ye et al., 2020) target specific aspects of call language. Of these, only a minority release code that still runs; §4 audits which.

The shared substrate. Nearly every paper above evaluates on label and split files derived from Qin and Yang (2019) and redistributed by successors (VolTAGE, Ke-FVP), rather than recomputing targets from prices. Yu et al. (2024) are the first to report that training-free baselines built on a company’s own past volatility beat every transcript model on these files, and release augmented earnings histories that make such baselines easy to build. No prior work audits the files themselves—embargo, train/test company overlap, label integrity, or the horizon convention—which is the gap §3.1 fills.

Audio and prosody. Chen et al. (2025) report that physics-informed acoustic features explain up to 43.8% of 30-day realised-volatility variance, motivating an audio ladder; the FinAudio benchmark (Cao et al., 2025) cautions that audio-LLMs struggle on long financial audio, informing our preference for frozen pooled features over end-to-end audio-LLMs in v1.

Known pitfalls (the methodological core). Our design is organised around four documented failure modes.

(1) *Ticker-identity leakage* (Yu et al., 2024) is the central threat to validity (§1); LLMs also deanonymise entity-neutered call transcripts with near-perfect accuracy (Lopez-Lira et al., 2025), so identity leaks even under masking. We respond with identity-robust targets (Δv , HAR-residual), a ticker-fixed-effect baseline, and an identity-control suite (§5). (2) *LLM lookahead bias*: a 2024-cutoff model “predicting” 2019–2021 outcomes it has effectively seen. A now-active literature offers diagnostics (Gao et al., 2025), cutoff-bounded pretraining (Yan et al., 2026; He et al., 2025), workflow benchmarks that find significant bias in open model families we use (Benhenda, 2026), and memorization filtering (Roy and Roy, 2026); the nearest precedent validates an LLM disagreement measure post-cutoff on ECB press conferences (Collodel, 2025). None evaluates a frozen earnings-call pipeline on post-cutoff calls; we respond with exactly that. (3) *Weak baselines*: much prior work omits persistence, GARCH(1,1), and HAR-RV (Corsi, 2009), which are hard to beat at long horizons; we compute these first and use persistence as the denominator of our headline metric (Gu et al., 2020). (4) *Gender bias in audio features* (Wang et al., 2024; Lin et al., 2024; Mathur et al., 2022): we analyse and report a gender confound rather than debias. We deliberately exclude agentic trading systems and time-series foundation models from v1 scope.

3 The ECVOL Benchmark

We first document the substrate the literature shares (§3.1), then describe what ECVOL-bench supplies in its place: corpora with reconstructed identity (§3.2), targets under both horizon conventions (§3.3), measured call timing (§3.4), leakage-proof splits (§3.5), and the release (§3.6).

3.1 The substrate the literature shares

The EarningsCall (EC) benchmark of Qin and Yang (2019) and the two MAEC releases (2015 and 2016 cohorts; Li et al., 2020) are evaluated, across the lineage in §2, on split and label files that are redistributed from paper to paper rather than recomputed. We fetched the canonical files from the VoltAGE, KeFVP, and HTML repositories, manifested them by SHA-256, and audited four properties. *[pending: Table 1 (Result Table 5R) is regenerated by ecvol audit substrate; the figures below were established in the 2026-08-23 audit and are quoted pending that regeneration.]* (1) *Provenance*. The VoltAGE and KeFVP test-split files are byte-identical: the lineage shares one physical label set, never recomputed from prices. (2) *Embargo*. Every train/validation and validation/test boundary has a zero-day gap—on EC, training ends and validation starts on the same date (2017-08-03; 4 training and 14 validation calls on that day)—against targets that run 30 days past each call, so every boundary leaks the target window

[pending: Result Table 5R: per benchmark (EC, MAEC-15, MAEC-16) and for our own splits as the contrast row—file SHA-256, boundary embargo in days, test-call company overlap, degenerate label cells, and horizon convention; each cell with the file and row counts it was computed from.]

Table 1: Substrate audit of the shared benchmark files. Regenerated by `ecvol audit substrate`.

(3) *Company overlap*. On the temporal splits the field reports, 80.4% (EC), 44.2% (MAEC-15), and 55.7% (MAEC-16) of test calls come from companies that also appear in training. A temporal split *always* permits this; the finding is not that overlap exists but that no paper in the lineage reports the disjoint condition, so the identity confound of Yu et al. (2024) has never been measured on these benchmarks. (4) *Label integrity and convention*. The single-day log-vol series contain 137 exact-zero cells out of 16,800 (0.82%; the headline averaged series is clean), and KeFVP’s MAEC files interleave four binary labels at exactly $\tau \in \{3, 7, 15, 30\}$ among 26 price columns under identical `future_label_N` naming. Finally, the source paper defines its horizons over *calendar* days; most downstream work, and our own design document until its correction, describes them as trading days—at $\tau=30$ the difference is roughly 21 versus 30 sessions, and every published number depends on which one was used.

3.2 Corpora and identity reconstruction

We build on two openly licensed multimodal earnings-call corpora and reserve a third, post-cutoff corpus for the lookahead and clean-audio experiments. **FinCall-Surprise** (Shu et al., 2025) (primary; Apache-2.0) contains 2,688 calls spanning 2019–2021 with full MP3 audio, JSON transcripts, and slides. **MAEC** (Li et al., 2020) (secondary; CC-BY-SA-4.0) contains $\sim 3,443$ calls over 2015–2018 (S&P 1500) with sentence-aligned audio features and transcripts but no raw audio. **Earnings25** [TODO: cite; Zenodo 10.5281/zenodo.18762168] (CC-BY-4.0; ~ 500 Q4-2025 S&P 500 calls, 498 h of audio, held January–February 2026) is post-cutoff for every model we use and carries native audio at source quality; it is the Phase-7 corpus (§5) *[pending: ingestion gated on a metadata check]*.

FinCall-Surprise ships with four fields per record and *no* ticker, company, CIK, date, time, fiscal period, or speaker names. We reconstruct identity and date from transcript content and slide metadata, cross-check against price data, and release the resulting table together with an 80-row manual override list: the reconstruction resolves 2,496/2,688 calls (92.9% overall; 95.3% of the 2,499 earnings-type calls) and passes a 60/60 manual audit. This table is what makes the identity controls of §5 runnable on the corpus at all.

3.3 Price data and targets

Daily adjusted-close OHLCV is obtained from a free, no-account source and cross-checked against Tiingo on a random ticker sample (adjusted-return correlation > 0.999); delisted or unmatched tickers produce an explicit coverage report rather than being silently dropped. After price joining, FinCall-Surprise yields 2,424 calls with at least one target (97.1% join rate) and MAEC yields 2,578 (75.4%; a documented shortfall driven by 2015–2018 delistings), for a combined modelling set of 5,002 calls.

For each call we define returns $r_t = (P_t - P_{t-1})/P_{t-1}$ relative to day $t = 0$, the last trading day before the call’s information is public (§3.4). The primary target is log realised volatility over horizon $\tau \in \{3, 7, 15, 30\}$,

$$v_{\text{post}}(\tau) = \ln\left(\sqrt{\frac{1}{\tau} \sum_{t=1}^{\tau} (r_t - \bar{r})^2}\right), \quad (1)$$

with pre-call volatility $v_{\text{pre}}(\tau)$ defined over the τ sessions ending at day 0. We ship every target under *both* horizon conventions: the trading-day convention (sessions $1..\tau$; all results in §6 use it) and the calendar-day convention of [Qin and Yang \(2019\)](#) that the published leaderboard inherits; every result table carries the convention as a column *[pending: calendar-day targets and the between-convention delta, T9.1]*. We add two identity-robust variants: the **change** $\Delta v(\tau) = v_{\text{post}}(\tau) - v_{\text{pre}}(\tau)$, which subtracts the company’s own level, and the **HAR-residual** $v_{\text{post}}(\tau) - \widehat{\text{HAR}}(\tau)$, deviation from a train-only HAR-RV ([Corsi, 2009](#)) forecast. Insufficient post-call history and zero-variance windows yield NaN targets excluded with a reason code; targets are unit-tested against hand-computed values.

3.4 Call timing

Fewer than 4% of FinCall-Surprise transcripts mention a clock time, and 32.4% of recovered *dates* come from slide-PDF creation stamps, so no call in either corpus carries a measured time of day. All results in §6 therefore use a uniform, flagged after-hours rule (calls assumed to occur after market close; $t=1$ is the next trading session). *[pending: T9.2 retrofits measured call datetimes with a provenance tier—first from the before/after-market attribute shipped with [Yu et al. \(2024\)](#)’s earnings histories, then from EDGAR 8-K Item 2.02 filings for the remainder—and reports the sensitivity of every target to the after-hours rule.]*

3.5 Leakage-proof splits

Three split families are generated once as committed CSVs and guarded by leakage assertions in CI. *Temporal*: train $<$ validation $<$ test in strict calendar order ($\sim 70/10/20$) with a **30-trading-day embargo** between segments so that no 30-day target window crosses a boundary—the constraint the shared substrate lacks

(§3.1). *Ticker-disjoint*: companies in validation/test never appear in training (grouped, seeded), the headline identity-control split; by construction its test-call company overlap is 0.0%, against 91.3% (FinCall-Surprise) and 78.9% (MAEC) on our own temporal splits. *Combined*: both constraints simultaneously, reported as the hardest robustness row. An information rule is codified and enforced: every feature function receives an `as_of` timestamp, and the harness fails any feature that reads price data after it. FinCall-Surprise spans COVID; per-year breakdowns are mandatory and a Feb–Apr 2020 exclusion is reported as an exploratory sensitivity.

3.6 Release

ECVOL-bench is released in full: the identity/date reconstruction table and override list, split CSVs, targets under both conventions, the derived-feature archives of §4, the evaluation harness and baselines, the substrate-audit report, and every result table. FinCall-Surprise-derived artifacts are Apache-2.0; MAEC-derived artifacts inherit CC-BY-SA-4.0. Raw audio and transcripts from sources with unclear licenses are never redistributed (a *scripts-not-data* policy).

4 Models

We organise our own models as a ladder (Stages 0–4; Figure 1). Two invariants make the comparison honest. First, from Stage 2 upward every learned model includes past-volatility covariates (v_{pre} at all horizons plus the HAR-RV forecast) and is *also* reported with them ablated, baking the persistence control into every comparison. Second, given the small-data regime (10^3 – 10^4 samples) we use frozen representations with light heads (ridge, shallow MLP, LightGBM) and per-horizon models by default; fine-tuning is gated to a conditional later stage and a shared multi-task head is an ablation, not the headline. Because the shallow-MLP heads proved seed-unstable (across-seed R_{OOS}^2 std up to ≈ 3.4 ; see Limitations), the consolidated grid (Table 6) uses one pre-registered deterministic ridge/structural canonical model per stage; MLP heads appear in the appendix only.

Stage 0 — econometric baselines. Persistence ($\hat{v} = v_{\text{pre}}$), EWMA (RiskMetrics $\lambda=0.94$), HAR-RV (OLS on daily/weekly/monthly realised-variance factors, [Corsi, 2009](#)), and GARCH(1,1) (per-call fit, `arch`). These establish the floor and the denominator of the headline metric.

Stage 1 — tabular identity baseline. LightGBM over past-volatility features, trivial metadata (transcript length, Q&A turn count where available), and a **ticker fixed effect**. This is the in-house “Same-Company-




Figure 1: The modelling ladder. From Stage 2 up, every model also carries past-volatility covariates (and is reported with them ablated). Every stage runs through the identity-control suite. Green = complete, orange = pending.

Same-Signal” control: it measures how much apparent multimodal performance is identity plus metadata.

Stage 2 — frozen text content. We segment each transcript into prepared-remarks vs. Q&A and speaker turns, then extract three feature blocks: BGE-M3 (BAAI, 2024) section-pooled sentence embeddings (1024-d), FinBERT (ProsusAI, 2019) sentiment aggregates per scope×role, and surface statistics. Heads are ridge (with the α grid selected on validation, *no* retrain on train+val) and a shallow MLP (1×256, L2 + early stopping, 5 seeds). Features use fp32 with content-hash caching for bit-identical re-extraction.

Stage 3 — frozen audio content. Audio is QC’d and resampled to 16 kHz mono FLAC. We extract eGeMAPSv02 functionals (88-d, openSMILE (Eyben et al., 2010), CPU, interpretable), then WavLM-Large (Chen et al., 2021b) and emotion2vec+ (Ma et al., 2023) embeddings over 30 s non-overlapping windows, mean-pooled to a per-call vector (2,671 calls, all extracted on one consumer GPU with resumable content-hash caches), fed to the same heads. We use per-call pooling only (no per-sentence alignment), and report a gender-confound analysis.

Stage 4 — fusion. Gated fusion (per-modality train-fit PCA + norm gates → shallow MLP) and late-fusion stacking (a meta-ridge over the text, audio, and Stage-1 GBDT base predictions, fit on validation only) over frozen modality embeddings, to test whether multimodal beats the best unimodal model (RQ2). Cross-attention is omitted: with one pooled vector per modality, attention over 2–3 tokens reduces to gated fusion. All fusion heads are <5M parameters.

Reproduced prior models. Of ~20 papers in the lineage of §2, five release code that still runs: HTML (Yang et al., 2020), Same-Company-Same-Signal (Yu

et al., 2024), KeFVP (Niu et al., 2023), DialogueGAT (Sang and Bao, 2022), and VoltAGE (Sawhney et al., 2020). We re-run each on its own original dataset, first reproducing the published-split number within a stated tolerance, then evaluating the same model under an embargoed and a ticker-disjoint split of that dataset, so that “the model fails” cannot be confused with “the model does not transfer” (RQ3). The remaining fifteen—among them the strongest recent positive claims (Cao et al., 2024a; Chen et al., 2025)—have no code, empty or unpushed repositories, closed-API dependencies, or a stated inability to run; we tabulate each with the evidence [pending: Phase 6R; Table 7 and the code-availability audit table in the appendix].

5 Experimental Protocol

Metrics. We report MSE on log-RV per horizon (literature comparability); R^2_{OOS} vs. persistence (Gu et al., 2020; $1 - \sum(y - \hat{y})^2 / \sum(y - \hat{y}_{\text{persist}})^2$), the headline metric, which is not gameable by identity memorisation the way raw MSE is; MAE; and cross-sectional Spearman rank correlation per calendar quarter. All metrics are reported on all three targets (v , Δv , HAR-residual) × all splits, and carry the horizon convention as a column.

Statistical rigor. Every learned model is run with ≥ 5 seeds (mean $\pm$ std). We use Diebold–Mariano tests (Diebold and Mariano, 1995) (HLN small-sample correction) on per-call squared-error differentials vs. persistence and vs. the best preceding stage; cluster-bootstrap confidence intervals (1,000 resamples) clustered by ticker and by calendar quarter; and Holm-corrected p -values across horizons for confirmatory claims. Confirmatory vs. exploratory analyses are labelled in advance: confirmatory analyses are Stage- k vs. Stage-1 DM tests on Δv at each τ , Stage 4 vs. best unimodal, the temporal-vs-ticker-disjoint gap per stage, and—for each reproduced prior model—the published-split vs. leakage-proof-split gap; all per-year/regime, COVID-exclusion, prepared-vs-Q&A, call-length, and gender-confound analyses are exploratory.

Identity-control suite. The signature diagnostic of this benchmark probes whether content models read content or identity, via four controls. (1) *Ticker-only model*: predict the per-ticker target mean (global mean for unseen tickers); if a content model matches it on level- v , that model reads identity. (2) *Same-ticker transcript shuffle*: replace each test call’s transcript with a different call from the *same* ticker; if performance is unchanged the model reads identity, not content, whereas a *global shuffle* (any other call) isolates pure overfitting. (3) *Identity probe*: a linear probe predicting the ticker from frozen call embeddings; high accuracy demonstrates identity is encoded. (4) *Ticker-disjoint split* re-

sults reported beside temporal-split results in every table. The suite runs unchanged on the reproduced prior models (§4).

Post-cutoff experiments. Earnings25 (§3.2) supports two experiments that the 2019–2021 corpora cannot. *Clean audio*: the frozen Stage-3 audio ladder (same extractors, same heads, same controls) is re-run on source-quality 2025 recordings, to test whether the audio-inert finding of §6 is a property of the task or of FinCall-Surprise’s uniform 40 kbps transcodes. *Lookahead*: the entire frozen pipeline—no retraining, no post-hoc threshold changes (rule fixed in advance)—is evaluated on calls held after the knowledge cutoff of every model used; disproportionate degradation relative to the 2019–2021 results would be contamination evidence. *[pending: Both are owed: T9.4 and Phase 7.]*

Reproducibility. All result tables regenerate byte-identically from committed CSV artifacts via `ecvol` report, guarded by a CI test. Every result CSV carries a committed run manifest—the resolved command config and its hash, git SHA, seeds, environment fingerprint, and the file’s SHA-256—and both `ecvol` report and CI refuse a result file whose bytes match no manifest; one committed config per result-producing command is the record of how it was run. Every data file carries a SHA-256 manifest entry; targets are unit-tested against hand-computed values; and the splits are guarded by leakage assertions in CI. Manifests written for results that predate this mechanism are labelled *backfill* until their producing command is re-run. *[pending: Clean-machine regeneration of Table 2 is recorded with the reproducibility package.]*

6 Results

We report results factually and defer interpretation to §7. Curated headline tables appear here on the identity-robust Δv target for FinCall-Surprise under the trading-day convention; the complete grid (both corpora, level- v and Δv , R_{OOS}^2 and MSE, all splits and horizons) is in Appendix A. Stars denote Diebold–Mariano significance ($p < 0.05$): in Table 2 vs. persistence, and in Tables 3 and 5 vs. the Stage-1 ticker-FE GBDT; in the consolidated grid (Table 6) markers are Holm-corrected, across horizons and signed (* better / † worse than Stage-1).

6.1 Baselines (Stages 0–1)

Table 2 shows the econometric floor. HAR-RV is strong at short and medium horizons on both splits, confirming RQ1’s premise that past-volatility baselines are hard to beat. The one striking failure is the FinCall *temporal*-split at $\tau=30$, where every model falls below persistence (HAR-RV $R_{OOS}^2 = -0.287$): the training win-

Model	$\tau=3$	$\tau=7$	$\tau=15$	$\tau=30$
<i>Temporal split (n=481)</i>				
Persistence	0.000	0.000	0.000	0.000
EWMA	0.387*	0.362*	0.145*	-0.219*
HAR-RV	0.464*	0.410*	0.207*	-0.287*
GARCH(1,1)	0.265*	0.193*	-0.140*	-0.541*
GBDT (ticker FE)	0.344*	0.248*	-0.057	-0.904*
<i>Ticker-disjoint split (n=484)</i>				
Persistence	0.000	0.000	0.000	0.000
EWMA	0.357*	0.199*	-0.009	0.046
HAR-RV	0.495*	0.357*	0.211*	0.229*
GARCH(1,1)	0.356*	0.229*	0.088	0.209*
GBDT (ticker FE)	0.407*	0.275*	0.089	0.108

Table 2: Baselines. R_{OOS}^2 vs. persistence on the Δv target, FinCall-Surprise (test). HAR-RV is strong except on the COVID-regime temporal $\tau=30$ cell. *: Diebold–Mariano significant vs. persistence ($p < 0.05$).

dow is COVID-inflated (Feb–May 2020) while the test window is calm late 2021, so a long-horizon level forecast trained on the crisis regime is mis-calibrated. The same models pass comfortably on the ticker-disjoint split (HAR-RV $R_{OOS}^2 = +0.229$ at $\tau=30$) and on MAEC, so we treat this as a documented regime exception rather than a target bug; per-year breakdowns are mandatory in the appendix.

6.2 Text content (Stage 2)

Table 3 shows frozen-text heads against the Stage-1 baseline on Δv . The result is a clean negative: *text content does not cleanly beat the identity/past-volatility baselines*. On the FinCall temporal split the text ridge head is near zero and collapses at $\tau=30$; the MLP heads overfit badly. On the ticker-disjoint split the text ridge head is weak (R_{OOS}^2 between -0.04 and $+0.17$) while the past-vol-only head is consistently strong ($R_{OOS}^2 \approx 0.19$ – 0.45), so any apparent text gain is not separable from what past volatility already provides. MAEC shows the same pattern at smaller magnitude (Appendix A).

6.3 Identity controls (Stage 2 diagnostic)

Table 4 explains the negative result. A linear probe recovers the ticker from frozen call embeddings at 89.5% on FinCall ($319\times$ chance) and 53.9% on MAEC ($487\times$ chance): identity is heavily encoded. The transcript shuffle is decisive—replacing each test call’s transcript with a *different call from the same ticker* barely changes R_{OOS}^2 , whereas a global shuffle degrades it. The heads are reading *which company* is speaking, not *what* is said. This is the in-corpus confirmation of Yu et al. (2024) and the reason RQ4 is central—and it is only measurable because the benchmark reconstructed the identity the corpus omits (§3.2).

Model	$\tau=3$	$\tau=7$	$\tau=15$	$\tau=30$
<i>Temporal split (n=481)</i>				
Ridge (text)	0.028*	0.030*	0.015	-0.144*
Ridge (past-vol)	0.415*	-1.053	-2.858	-0.667
Ridge (text+vol)	-1.678	-3.076	-3.586	-0.468
MLP (text)	-0.108*	-0.175*	-0.265*	-1.042
<i>Ticker-disjoint split (n=484)</i>				
Ridge (text)	-0.044*	-0.026*	0.008	0.171
Ridge (past-vol)	0.451	0.306	0.191*	0.224*
Ridge (text+vol)	0.223*	0.129*	0.094	0.236
MLP (past-vol)	0.492*	0.363*	0.230*	0.249*

Table 3: Text content (Stage 2). R^2_{OOS} vs. persistence on the Δv target, FinCall-Surprise (test). The text-only head does not separate from the past-volatility head; adding text to past-vol does not help. *: Diebold–Mariano significant vs. the Stage-1 ticker-FE GBDT ($p < 0.05$). Full head set in Appendix A.

Corpus	Identity probe (linear)	
	Accuracy	$\times$ chance
FinCall ($n=2656$, 357 tickers)	89.5%	319 $\times$
MAEC ($n=3132$, 902 tickers)	53.9%	487 $\times$

Condition	$\tau=3$	$\tau=7$	$\tau=15$	$\tau=30$
Real transcript	0.028	0.030	0.015	-0.144
Within-ticker shuffle	0.017	0.020	0.006	-0.142
Global shuffle	-0.099	-0.121	-0.115	-0.266

Table 4: Identity controls. Top: a linear probe predicts the ticker from frozen BGE-M3 call embeddings far above chance, demonstrating identity is heavily encoded. Bottom: R^2_{OOS} of the text ridge head on the Δv target (FinCall, temporal split, test, $n=481$). Replacing each test call’s transcript with a different call from the same ticker (within-ticker shuffle) barely changes performance, whereas a global shuffle degrades it—evidence the head reads ticker identity rather than call content.

6.4 Audio content (Stage 3)

Table 5 extends the same design to audio. Audio-only heads (eGeMAPS, WavLM-Large, emotion2vec+, ridge or MLP) sit at R^2_{OOS} zero on every split and horizon (−0.06 to +0.03): frozen audio content alone carries no measurable Δv signal. Adding past-volatility covariates reproduces the Stage-2 pattern—the only DM-significant wins over Stage-1 are past-vol-augmented heads on the ticker-disjoint split. The audio-shuffle control (bottom block) is decisive: replacing each test call’s audio with a random other company’s call leaves most of those gains intact ($\tau=3$: 0.337 $\rightarrow$ 0.318; $\tau=30$: 0.212 $\rightarrow$ 0.100), so the audio features are inert and past volatility carries the signal. Identity is nonetheless encoded: a linear probe recovers the ticker from per-call WavLM embeddings at 76.3% (271 $\times$ chance; emotion2vec+ 29.4%, 105 $\times$). An exploratory gender-confound check (eGeMAPS-F0 dominant-speaker proxy, 100% coverage) finds near-identical per-group test MSE (0.241 low-pitch vs. 0.242

Model	$\tau=3$	$\tau=7$	$\tau=15$	$\tau=30$
<i>Temporal split (n=481)</i>				
Ridge (eGeMAPS)	-0.000*	-0.007*	-0.015	-0.007*
Ridge (WavLM)	-0.005*	-0.007*	-0.045	-0.031*
Ridge (emotion2vec+)	-0.018*	-0.028*	-0.056	-0.060*
Ridge (WavLM+past-vol)	0.265	-1.026	-2.032	-0.561
<i>Ticker-disjoint split (n=484)</i>				
Ridge (eGeMAPS)	-0.011*	-0.013*	-0.002	0.002
Ridge (WavLM)	-0.038*	-0.035*	-0.027*	0.024
Ridge (emotion2vec+)	-0.002*	-0.003*	0.011	0.002
Ridge (WavLM+past-vol)	0.337*	0.171*	0.111	0.212*
<i>Audio-shuffle control of WavLM+past-vol (ticker-disjoint)</i>				
Real audio	0.337	0.171	0.111	0.212
Within-ticker shuffle	0.317	0.144	0.097	0.170
Global shuffle	0.318	0.115	0.093	0.100

Table 5: Audio content (Stage 3). R^2_{OOS} vs. persistence on the Δv target, FinCall-Surprise (test). Audio-only heads sit at zero everywhere; the apparent gains of the past-vol-augmented head survive replacing every test call’s audio with a random other company’s call (global shuffle, bottom block), so past volatility, not audio content, carries them. *: Diebold–Mariano significant vs. the Stage-1 ticker-FE GBDT ($p < 0.05$; read direction from the value). Full head set in Appendix A.

high-pitch) and a 0.03 correlation between speaker F0 and squared error: with inert audio features, no large gender disparity is detectable (see the Ethics Statement for caveats). All of this rests on FinCall-Surprise’s 40 kbps source audio; §6.8 is the test of whether it transfers.

6.5 Fusion and the consolidated grid (Stage 4)

Gated fusion and late-fusion stacking over all frozen modalities answer RQ2 in the negative: *multimodal is not better than the best unimodal model*. Only one of 48 Δv test cells beats Stages 1–3 DM-significantly (gated fusion, temporal $\tau=7$, $R^2_{\text{OOS}} = +0.019$ —negligible), and the best fusion cell overall (stacking, ticker-disjoint $\tau=3$, $R^2_{\text{OOS}}=0.436$) is not significant vs. Stage-1 ($p=0.13$). Table 6 consolidates the ladder with one pre-registered canonical model per stage and Holm-corrected confirmatory tests. Two facts summarise it. On the regime-honest *temporal* split, no content stage beats the econometric floor at any horizon; where a content model is Holm-significantly better than Stage-1 (fusion at $\tau=30$) it merely beats a collapsed GBDT while remaining below persistence. On the *ticker-disjoint* split the favorable cells are exactly the past-vol-augmented ones that the shuffle controls (Tables 4 and 5) show to be identity/past-vol-confounded. The per-year breakdown sharpens the regime point: the FinCall temporal test year is entirely 2021, and there every content model underperforms persistence on level- v (HAR −0.29, GBDT −0.82, text −1.09, audio −1.66, fusion −1.02).

Stage (canonical model)	$\tau=3$	$\tau=7$	$\tau=15$	$\tau=30$
<i>Temporal split (n=481)</i>				
S0 HAR-RV	0.464	0.410	0.207	-0.28*
S1 GBDT (ticker FE)	0.344	0.248	-0.057	-0.90†
S2 Text (ridge+vol)	-1.678	-3.076	-3.586	-0.468
S3 Audio (ridge+vol)	0.214	-1.099	-3.706	-0.548
S4 Fusion (stack)	0.232	-0.326	-1.828	-0.106*
<i>Ticker-disjoint split (n=484)</i>				
S0 HAR-RV	0.495	0.357	0.211	-0.229*
S1 GBDT (ticker FE)	0.407	0.275	0.089	-0.108*
S2 Text (ridge+vol)	0.223†	0.129†	0.094	-0.236†
S3 Audio (ridge+vol)	0.304†	0.147†	0.044	0.209
S4 Fusion (stack)	0.436	0.292	0.133	0.244*

Table 6: Consolidated grid (Stages 0–4), pre-registered canonical model per stage. R_{OOS}^2 vs. persistence on the Δv target, FinCall-Surprise (test); persistence is 0 by construction. Significance is the Holm-corrected confirmatory family (Stage- k vs. Stage-1 DM across horizons): * = significantly better than Stage-1, † = significantly worse ($p < 0.05$). No content stage beats the econometric floor on the temporal split (the fusion $\tau=30$ star beats only the collapsed GBDT, not persistence); the ticker-disjoint gains are past-vol/identity-confounded (Table 5, bottom block). Canonical models are deterministic ridge/structural heads; seed-unstable MLP heads are reported in Appendix A only.

[pending: Result Table 6R.]

Table 7: Reproduced prior models under leakage-proof evaluation. Each model on its own original dataset; published vs. reproduced vs. embargoed vs. ticker-disjoint.

6.6 Horizon convention: calendar vs. trading days

[pending: T9.1: every table above recomputed under the calendar-day convention of Qin and Yang (2019), the between-convention delta per stage, split, and horizon, reported as a finding about how much the published leaderboard depends on an undocumented convention.]

6.7 Prior models under the new controls (Stage R)

[pending: Phase 6R, Result Table 6R (Table 7): for each of HTML, Same-Company-Same-Signal, KeFVP, DialogueGAT, and VolTAGE on its own dataset—the published-split number as reported, our reproduction of it with tolerance, and the same model under an embargoed and a ticker-disjoint split, with the identity probe and shuffle controls attached. Every model gets either both split conditions or a documented reason code.]

6.8 Clean audio: the Stage-3 ladder on Earnings25

[pending: T9.4: the frozen eGeMAPS / WavLM / emotion2vec+ ladder with the same heads and controls as

Table 5, on source-quality 2025 audio; real vs. shuffle reported; any divergence from the FinCall result stated with the audio-quality difference as the candidate cause.]

6.9 Post-cutoff lookahead

[pending: Phase 7: the frozen-pipeline evaluation on Earnings25’s post-cutoff calls, no retraining, rule fixed in advance.]

7 Discussion

What the evidence shows. On the completed ladder (Stages 0–4), four statements hold. (i) Strong econometric baselines are hard to beat, and the headline metric (R_{OOS}^2 vs. persistence) exposes a long-horizon regime failure (FinCall temporal $\tau=30$) that raw MSE would obscure. (ii) Frozen text content does not cleanly add signal over past-volatility and identity baselines on the identity-robust Δv target, and the reason is measurable: call embeddings encode ticker identity (89.5% probe accuracy on FinCall), and within-ticker transcript shuffles leave performance essentially unchanged—an in-corpus replication of the “Same Company, Same Signal” effect (Yu et al., 2024). (iii) Frozen audio content is inert: audio-only heads sit at R_{OOS}^2 zero everywhere, the apparent gains of past-vol-augmented audio heads survive a global audio shuffle, and WavLM embeddings still encode the speaker’s company at 76.3% probe accuracy—identity without predictive content. (iv) Fusion does not rescue the modalities it fuses: multimodal never significantly beats the best unimodal model (RQ2 negative). Together these form one coherent finding: in this task family, what looks like content signal is overwhelmingly company identity plus past volatility.

Why this is a benchmark paper. Every one of those statements depends on apparatus the shared substrate does not have. The identity probe and the shuffle controls need a ticker per call, which FinCall-Surprise does not ship; the ticker-disjoint rows need a split the lineage never reports; the regime finding needs per-year breakdowns over recomputed targets rather than inherited label files; and the embargo makes the temporal numbers honest at all (§3.1). A field that evaluates on files with a zero-day embargo, majority company overlap, and an unstated horizon convention cannot distinguish a model from a memory of which company it is looking at. ECVOL-bench is the repair, and the negative results above are what the repaired substrate reports for the modern representations we could run. The reproduction study (§6.7) is the same test applied to the field’s own models on the field’s own data; it is the part of the argument this draft still owes.

Engaging the strongest positive claims. The strongest recent positive claim—up to 43.8% of 30-day realised-volatility variance from text+acoustic features (Chen et al., 2025)—reports no identity controls, no released code, and a proprietary corpus. Prosody was the modality most plausibly carrying content orthogonal to identity; our modern audio ladder (WavLM, emotion2vec+, eGeMAPS) finds it inert once past volatility is controlled—albeit on the corpus’s 40 kbps source audio (see Limitations and §6.8)—and our shuffle controls predict such gains would not survive a within-ticker (let alone global) shuffle, a check the released suite makes cheap. Likewise, ECC Analyzer (Cao et al., 2024a) reports large MSE reductions from LLM-extracted semantics with closed-API components and no released code; it cannot be reproduced, and we say so in the code-availability audit rather than argue with it.

The pre-registered framing gate. The study was pre-registered with a framing gate: a “positive showcase” spine was to be adopted only if content-bearing models beat the Stage-0/1 floor on Δv or HAR-residual with DM-significant improvements that survived the identity-control suite on at least two of four horizons. The criterion was tested and missed at Stage 2 (text), Stage 3 (audio), and Stage 4 (fusion), and the gate was recorded as resolved toward re-examination in the project’s decision log before the benchmark framing was adopted. We report this so that the framing of this paper can be audited against what the data said and when.

What a clean answer would need. Independent of any model, two facts about the corpora limit every result in this literature and are addressed only by new data: the uniform 40 kbps transcode underlying all FinCall Surprise acoustics, and the absence of any measured call time of day. The first is why the Earnings25 clean audio re-run (§6.8) is the single most informative pending experiment—it either makes the audio-inert finding near-unassailable or overturns it; the second is why the timestamp retrofit (§3.4) is part of the benchmark rather than a footnote. **[TODO: Finalise this section once Phase 6R, T9.1, T9.4, and Phase 7 land.]**

8 Conclusion

We audited the evaluation substrate that a decade of multimodal earnings-call volatility prediction shares— inherited label files with a zero-day embargo against 30-day targets, majority company overlap between training and test, no reported ticker-disjoint condition, degenerate label cells, and a horizon convention that down-stream work misattributes—and released ECVOL-bench to replace it: reconstructed identity and dates for the one corpus with open audio, targets under both horizon conventions, embargoed and ticker-disjoint splits, strong econometric baselines, and an identity-control

suite, all regenerable from one command line. Run over modern open-weight text and audio representations, the benchmark’s answer is consistent and negative: frozen text content reads company identity rather than call content (89.5% ticker-probe accuracy; within-ticker shuffles leave performance unchanged); frozen audio content is inert beyond past volatility (its apparent gains survive a global audio shuffle); and fusion never significantly beats the best unimodal model. What the leaderboard literature measures as content signal is, under honest controls, overwhelmingly company identity plus past volatility. *[pending: The reproduction study, the calendar-day recomputation, the clean-audio re-run, and the post-cutoff lookahead complete the argument.]* **[TODO: Finalise the verdict once the pending results land.]**

Limitations

Pending results. The substrate-audit numbers in §3.1 were established in a one-off audit and await regeneration by `ecvol audit substrate`; the reproduction study (§6.7), the calendar-day recomputation (§6.6), the clean-audio re-run (§6.8), and the post-cutoff lookahead (§6.9) are owed. Until the lookahead lands, results on 2019–2021 calls from models with later training cutoffs carry contamination risk by default. **Reproductions run on their own data.** Each prior model is re-evaluated on its original dataset, not on ECVOL-bench, to keep “the model fails” separable from “the model does not transfer”; two reproductions carry known costs (KeFVP’s MAEC embeddings are unavailable upstream; DialogueGAT stops at $\tau=15$ and requires a corpus rebuild). **MLP-head instability.** The shallow-MLP heads are seed-unstable (across-seed R^2_{OOS} std up to ≈ 3.4), which is why the consolidated grid uses deterministic ridge/structural canonicals; MLP cells in the appendix should be read as seed means with that caveat. **The past-vol covariate block is not HAR.** The learned $[v_{\text{pre}}, \text{RV}_{1/5/22}]$ ridge used inside Stage-2–4 heads overfits the COVID regime shift on the temporal split ($R^2_{\text{OOS}} - 2.0$ at $\tau=15$) where the rigid OLS HAR-RV stays robust (+0.21); the two are labelled distinctly throughout, and the gap is itself a small finding about structural vs. unconstrained conditioning under regime shift. **Short-horizon R^2_{OOS} inflation.** The persistence denominator is noisiest at $\tau=3$ (its MSE is $\sim 2.8 \times$ the $\tau=7$ one), so short-horizon R^2_{OOS} gains are easier to show and should be read accordingly. **Timestamps.** No call in either corpus carries a measured time of day, so every target rests on a flagged after-hours rule; misdated calls add noise to short-horizon targets, and the sensitivity of the results to that rule is unmeasured until the retrofit of §3.4 lands. **Regime sensitivity.** FinCall-Surprise spans COVID; the temporal $\tau=30$ baseline failure shows long-horizon level forecasts are regime-sensitive, which is why we emphasise the identity-robust Δv target and report per-year breakdowns. **MAEC cov-**

erage. The 75.4% price-join rate (2015–2018 delistings) makes MAEC a secondary, comparability corpus; it also ships no raw audio, so the audio ladder is FinCall-only. **Source audio quality.** Every FinCall audio file is a uniform 40 kbps MP3 transcode (eight source sample rates; 14 calls at telephone-grade 8 kHz), so aggressive lossy compression underlies all acoustic features. The audio-inert finding is therefore established on compression-damaged audio and may not transfer to studio-quality recordings—a caveat that applies equally to prior positive claims built on the same corpora, and the reason for the Earnings25 re-run. **No diarisation in v1.** Audio uses per-call pooling, not speaker-resolved features; the gender analysis correspondingly uses a dominant-speaker F0 proxy. **Identity reconstruction is imperfect.** 7.1% of FinCall-Surprise calls (4.7% of earnings calls) remain unresolved and are excluded with a reason code; the released override table records every manual decision.

Ethics Statement

Data and licensing. We use only openly licensed corpora (FinCall-Surprise, Apache-2.0; MAEC, CC-BY-SA-4.0) and open-weight models. Derived numeric features inherit and explicitly state their upstream license: raw audio/transcripts from unclear-license sources are never redistributed (*scripts-not-data*). **Gender confound.** Audio features encode speaker gender and female executives are underrepresented in these corpora, which can produce error disparities (Wang et al., 2024; Lin et al., 2024; Mathur et al., 2022); we analyse and report a gender-confound diagnostic rather than silently propagate the bias. **Intended use.** This is a forecasting-research benchmark, not investment advice; predictions are noisy and regime-sensitive, and should not be used for automated trading. **Reproducibility as an ethical commitment.** Given the documented reproducibility crisis in financial NLP (Tatarinov et al., 2025), we pin the environment, checksum every data file, and regenerate all result tables byte-identically from committed artifacts.

References

BAAI. BGE-M3 sentence embeddings. Hugging Face (MIT), 2024. URL <https://huggingface.co/BAAI/bge-m3>.

Mostapha Benhenda. Look-ahead-bench: a standardized benchmark of look-ahead bias in point-in-time LLMs for finance. arXiv:2601.13770, 2026. URL <https://arxiv.org/abs/2601.13770>.

Yupeng Cao et al. ECC analyzer: Extracting trading signal from earnings conference calls using large language models for stock volatility prediction. In *Proceedings of the 5th ACM International Conference*

on AI in Finance (ICAIF), 2024a. URL <https://arxiv.org/abs/2404.18470>.

Yupeng Cao et al. RiskLabs: Predicting financial risk using large language model based on multimodal and multi-sources data. arXiv:2404.07452, 2024b. URL <https://arxiv.org/abs/2404.07452>.

Yupeng Cao et al. FinAudio: A benchmark for audio large language models in financial applications. arXiv:2503.20990, 2025. URL <https://arxiv.org/abs/2503.20990>.

Chen et al. Distilling numeral information for volatility forecasting. In *Proceedings of the 30th ACM International Conference on Information and Knowledge Management (CIKM)*, 2021a. doi: 10.1145/3459637.3482089.

Sanyuan Chen et al. WavLM: Large-scale self-supervised pre-training for full stack speech processing. Hugging Face (MIT): microsoft/wavlm-large, 2021b. URL <https://huggingface.co/microsoft/wavlm-large>.

Xiaoliang Chen et al. The sound of risk: A multimodal physics-informed acoustic model for forecasting market volatility and enhancing market interpretability. arXiv:2508.18653, 2025. URL <https://arxiv.org/abs/2508.18653>.

Umberto Collodel. Interpreting the interpreter: Can we model post-ECB conferences volatility with LLM agents? arXiv:2508.13635, 2025. URL <https://arxiv.org/abs/2508.13635>.

Fulvio Corsi. A simple approximate long-memory model of realized volatility. *Journal of Financial Econometrics*, 2009. doi: 10.1093/jjfinec/nbp001.

Francis X. Diebold and Roberto S. Mariano. Comparing predictive accuracy. *Journal of Business & Economic Statistics*, 1995. doi: 10.1080/07350015.1995.10524599.

Florian Eyben, Martin Wöllmer, and Björn Schuller. openSMILE / eGeMAPS acoustic feature extraction. github.com/audeerling/opensmile-python, 2010. URL <https://github.com/audeerling/opensmile-python>.

Zhenyu Gao, Wenxi Jiang, and Yutong Yan. Detecting lookahead bias in LLM forecasts. arXiv:2512.23847, 2025. URL <https://arxiv.org/abs/2512.23847>.

Shihao Gu, Bryan Kelly, and Dacheng Xiu. Empirical asset pricing via machine learning. *The Review of Financial Studies*, 2020. doi: 10.1093/rfs/hhaa009.

Songrun He, Linying Lv, Asaf Manela, and Jimmy Wu. Chronologically consistent large language models. arXiv:2502.21206, 2025. URL <https://arxiv.org/abs/2502.21206>.

- Jiazheng Li, Linyi Yang, Barry Smyth, and Ruihai Dong. MAEC: A multimodal aligned earnings conference call dataset for financial risk prediction. In *Proceedings of the 29th ACM International Conference on Information and Knowledge Management (CIKM)*, 2020. doi: 10.1145/3340531.3412879.
- Yi-Cheng Lin et al. Emo-bias: A large scale evaluation of social bias on speech emotion recognition. arXiv:2406.05065, 2024. URL <https://arxiv.org/abs/2406.05065>.
- Y. Liu et al. AT-FinGPT: Financial risk prediction via an audio-text large language model. *Finance Research Letters*, 2025. URL <https://www.sciencedirect.com/science/article/abs/pii/S1544612325002314>.
- Alejandro Lopez-Lira, Yuehua Tang, and Mingyin Zhu. The memorization problem: Can we trust LLMs’ economic forecasts? arXiv:2504.14765, 2025. URL <https://arxiv.org/abs/2504.14765>.
- Ziyang Ma et al. emotion2vec: Self-supervised pre-training for speech emotion representation. Hugging Face: emotion2vec/emotion2vec_plus_large, 2023. URL https://huggingface.co/emotion2vec/emotion2vec_plus_large.
- Puneet Mathur et al. DocFin: Multimodal financial prediction and bias mitigation using semi-structured documents. In *Findings of the Association for Computational Linguistics: EMNLP 2022*, 2022. URL <https://aclanthology.org/2022.findings-emnlp.139/>.
- Hao Niu et al. KeFVP: Knowledge-enhanced financial volatility prediction. In *Findings of the Association for Computational Linguistics: EMNLP 2023*, 2023. URL <https://aclanthology.org/2023.findings-emnlp.770/>.
- ProsusAI. FinBERT: Financial sentiment analysis with bert. Hugging Face, 2019. URL <https://huggingface.co/ProsusAI/finbert>.
- Yu Qin and Yi Yang. What you say and how you say it matters: Predicting stock volatility using verbal and vocal cues. In *Proceedings of the 57th Annual Meeting of the Association for Computational Linguistics (ACL)*, 2019. URL <https://aclanthology.org/P19-1038/>.
- Anisha Roy and Dip Roy. MemGuard-Alpha: Detecting and filtering memorization-contaminated signals in LLM-based financial forecasting via membership inference and cross-model disagreement. arXiv:2603.26797, 2026. URL <https://arxiv.org/abs/2603.26797>.
- Yunxin Sang and Yang Bao. DialogueGAT: A graph attention network for financial risk prediction by modeling the dialogues in earnings conference calls. In *Findings of the Association for Computational Linguistics: EMNLP 2022*, 2022. URL <https://aclanthology.org/2022.findings-emnlp.117/>.
- Ramit Sawhney, Piyush Khanna, Arshiya Aggarwal, Taru Jain, Puneet Mathur, and Rajiv Ratn Shah. VolTAGE: Volatility forecasting via text audio fusion with graph convolution networks for earnings calls. In *Proceedings of the 2020 Conference on Empirical Methods in Natural Language Processing (EMNLP)*, 2020. URL <https://aclanthology.org/2020.emnlp-main.643/>.
- Shi et al. Enhancing volatility forecasting in financial markets: A general numeral attachment dataset for understanding earnings calls. In *Proceedings of IJCNLP-AACL 2023*, 2023. URL <https://aclanthology.org/2023.ijcnlp-short.5/>.
- Dong Shu, Yanguang Liu, Huopu Zhang, and Mengnan Du. FinCall-Surprise: A large scale multi-modal benchmark for earning surprise prediction. arXiv:2510.03965; data: github.com/Tizzzy/FinCall-Surprise (Apache-2.0), 2025. URL <https://arxiv.org/abs/2510.03965>.
- Nikita Tatarinov et al. Language modeling for the future of finance: A survey into metrics, tasks, and data opportunities. arXiv:2504.07274, 2025. URL <https://arxiv.org/abs/2504.07274>.
- Wang et al. AMA-LSTM: Pioneering robust and fair financial audio analysis for stock volatility prediction. NAACL 2024 Industry Track; arXiv:2407.18324, 2024. URL <https://arxiv.org/abs/2407.18324>.
- Yutong Yan et al. DatedGPT: Preventing lookahead bias in large language models with time-aware pretraining. arXiv:2603.11838, 2026. URL <https://arxiv.org/abs/2603.11838>.
- Linyi Yang, Tin Lok James Ng, Barry Smyth, and Ruihai Dong. HTML: Hierarchical transformer-based multi-task learning for volatility prediction. In *Proceedings of The Web Conference (WWW)*, 2020.
- Linyi Yang et al. NumHTML: Numeric-oriented hierarchical transformer model for multi-task financial forecasting. In *Proceedings of the AAAI Conference on Artificial Intelligence*, 2022. URL <https://arxiv.org/abs/2201.01770>.
- Ye, Qin, and Xu. Financial risk prediction with multi-round q&a attention network. In *Proceedings of the 29th International Joint Conference on Artificial*

Model	$\tau=3$	$\tau=7$	$\tau=15$	$\tau=30$
Persistence	0.000	0.000	0.000	0.000
EWMA	0.387*	0.362*	0.145*	-0.219*
HAR-RV	0.464*	0.410*	0.207*	-0.287*
GARCH(1,1)	0.265*	0.193*	-0.140*	-0.541*
GBDT (ticker FE)	0.354*	0.262*	-0.040	-0.839*

Table 8: fincall — temporal split — level- v — R^2_{OOS} vs persistence (test) (n=481).

Model	$\tau=3$	$\tau=7$	$\tau=15$	$\tau=30$
Persistence	0.000	0.000	0.000	0.000
EWMA	0.357*	0.199*	-0.009	0.046
HAR-RV	0.495*	0.357*	0.211*	0.229*
GARCH(1,1)	0.356*	0.229*	0.088	0.209*
GBDT (ticker FE)	0.412*	0.278*	0.078	0.116*

Table 9: fincall — ticker_disjoint split — level- v — R^2_{OOS} vs persistence (test) (n=484).

Table 10: fincall — temporal split — Δv — R^2_{OOS} vs persistence (test) (n=481).

Model	$\tau=3$	$\tau=7$	$\tau=15$	$\tau=30$
Persistence	0.000	0.000	0.000	0.000
EWMA	0.357*	0.199*	-0.009	0.046
HAR-RV	0.495*	0.357*	0.211*	0.229*
GARCH(1,1)	0.356*	0.229*	0.088	0.209*
GBDT (ticker FE)	0.407*	0.275*	0.089	0.108

Table 11: fincall — ticker_disjoint split — Δv — R^2_{OOS} vs persistence (test) (n=484).

Model	$\tau=3$	$\tau=7$	$\tau=15$	$\tau=30$
Persistence	0.000	0.000	0.000	0.000
EWMA	0.444*	0.301*	0.120*	-0.144*
HAR-RV	0.507*	0.392*	0.306*	0.206*
GARCH(1,1)	0.387*	0.281*	0.054	-0.309*
GBDT (ticker FE)	0.397*	0.333*	0.217*	0.131*

Table 12: maec — temporal split — level- v — R^2_{OOS} vs persistence (test) (n=507).

Model	$\tau=3$	$\tau=7$	$\tau=15$	$\tau=30$
Persistence	0.000	0.000	0.000	0.000
EWMA	0.398*	0.297*	0.119*	-0.144*
HAR-RV	0.474*	0.379*	0.311*	0.225*
GARCH(1,1)	0.306*	0.213*	0.009	-0.321*
GBDT (ticker FE)	0.395*	0.306*	0.236*	0.221*

Table 13: maec — ticker_disjoint split — level- v — R^2_{OOS} vs persistence (test) (n=514).

Model	$\tau=3$	$\tau=7$	$\tau=15$	$\tau=30$
Persistence	0.000	0.000	0.000	0.000
EWMA	0.444*	0.301*	0.120*	-0.144*
HAR-RV	0.507*	0.392*	0.306*	0.206*
GARCH(1,1)	0.387*	0.281*	0.054	-0.309*
GBDT (ticker FE)	0.407*	0.340*	0.213*	0.186*

Table 14: maec — temporal split — Δv — R^2_{OOS} vs persistence (test) (n=507).

Intelligence (IJCAI), 2020. URL <https://www.ijcai.org/proceedings/2020/631>.

Ding Yu et al. Same company, same signal: The role of identity in earnings call transcripts. arXiv:2412.18029, 2024. URL <https://arxiv.org/abs/2412.18029>.

A Complete Results

The tables below are generated byte-identically by `ecvol` report from the committed result CSVs (`data/results/`) and are the source of truth for the curated tables in §6; the curated tables copy selected cells from them. They span both corpora (FinCall-Surprise, MAEC), the level- v and Δv targets, both the R^2_{OOS} and MSE metrics, and the temporal and ticker-disjoint splits at all horizons, under the trading-day convention [pending: *calendar-day columns, T9.1*].

A.1 Stage-0/1 baselines (Result Table 1)

Model	$\tau=3$	$\tau=7$	$\tau=15$	$\tau=30$
Persistence	0.000	0.000	0.000	0.000
EWMA	0.398*	0.297*	0.119*	-0.144*
HAR-RV	0.474*	0.379*	0.311*	0.225*
GARCH(1,1)	0.306*	0.213*	0.009	-0.321*
GBDT (ticker FE)	0.394*	0.285*	0.227*	0.233*

Table 15: maec — ticker_disjoint split — Δv — R^2_{OOS} vs persistence (test) (n=514).

Model	$\tau=3$	$\tau=7$	$\tau=15$	$\tau=30$
Persistence	1.383	0.571	0.318	0.189
EWMA	0.768*	0.399*	0.280*	0.216*
HAR-RV	0.682*	0.347*	0.221*	0.150*
GARCH(1,1)	0.848*	0.411*	0.301	0.247*
GBDT (ticker FE)	0.834*	0.381*	0.249*	0.164*

Table 20: maec — temporal split — level-v — MSE (test) (n=507).

Model	$\tau=3$	$\tau=7$	$\tau=15$	$\tau=30$
Persistence	1.174	0.419	0.208	0.111
EWMA	0.720*	0.267*	0.178*	0.136*
HAR-RV	0.630*	0.247*	0.165*	0.143*
GARCH(1,1)	0.860*	0.335*	0.233*	0.167*
GBDT (ticker FE)	0.758*	0.309*	0.216	0.205*

Table 16: fincall — temporal split — level-v — MSE (test) (n=481).

Model	$\tau=3$	$\tau=7$	$\tau=15$	$\tau=30$
Persistence	1.268	0.492	0.302	0.184
EWMA	0.764*	0.346*	0.266*	0.210*
HAR-RV	0.667*	0.306*	0.208*	0.142*
GARCH(1,1)	0.881*	0.384*	0.296	0.242*
GBDT (ticker FE)	0.767*	0.342*	0.230*	0.143*

Model	$\tau=3$	$\tau=7$	$\tau=15$	$\tau=30$
Persistence	1.142	0.386	0.214	0.233
EWMA	0.735*	0.309*	0.216	0.222
HAR-RV	0.577*	0.248*	0.169*	0.180*
GARCH(1,1)	0.735*	0.297*	0.195	0.184*
GBDT (ticker FE)	0.672*	0.278*	0.197	0.206*

Table 17: fincall — ticker_disjoint split — level-v — MSE (test) (n=484).

Table 21: maec — ticker_disjoint split — level-v — MSE (test) (n=514).

Model	$\tau=3$	$\tau=7$	$\tau=15$	$\tau=30$
Persistence	1.383	0.571	0.318	0.189
EWMA	0.768*	0.399*	0.280*	0.216*
HAR-RV	0.682*	0.347*	0.221*	0.150*
GARCH(1,1)	0.848*	0.411*	0.301	0.247*
GBDT (ticker FE)	0.820*	0.377*	0.250*	0.154*

Model	$\tau=3$	$\tau=7$	$\tau=15$	$\tau=30$
Persistence	1.174	0.419	0.208	0.111
EWMA	0.720*	0.267*	0.178*	0.136*
HAR-RV	0.630*	0.247*	0.165*	0.143*
GARCH(1,1)	0.860*	0.335*	0.233*	0.167*
GBDT (ticker FE)	0.771*	0.315*	0.220	0.212*

Table 22: maec — temporal split — Δv — MSE (test) (n=507).

Table 18: fincall — temporal split — Δv — MSE (test) (n=481).

Model	$\tau=3$	$\tau=7$	$\tau=15$	$\tau=30$
Persistence	1.268	0.492	0.302	0.184
EWMA	0.764*	0.346*	0.266*	0.210*
HAR-RV	0.667*	0.306*	0.208*	0.142*
GARCH(1,1)	0.881*	0.384*	0.296	0.242*
GBDT (ticker FE)	0.769*	0.352*	0.233*	0.141*

Table 23: maec — ticker_disjoint split — Δv — MSE (test) (n=514).

Table 19: fincall — ticker_disjoint split — Δv — MSE (test) (n=484).

Model	$\tau=3$	$\tau=7$	$\tau=15$	$\tau=30$	Model	$\tau=3$	$\tau=7$	$\tau=15$	$\tau=30$
Ridge (text)	0.439*	0.319	0.037	-0.577	Ridge (text)	0.028*	0.030*	0.015	-0.144*
MLP (text)	0.147*	-0.241*	-1.101*	-2.633*	MLP (text)	-0.108*	-0.175*	-0.265*	-1.042
Ridge (past-vol)	0.415*	-1.052	-2.025	-0.371*	Ridge (past-vol)	0.415*	-1.053	-2.858	-0.667
MLP (past-vol)	0.456*	-0.228	-2.706	-2.477	MLP (past-vol)	0.390	-1.017	-2.249	-0.760
Ridge (text+vol)	0.448*	0.328	0.045	-1.091	Ridge (text+vol)	-1.678	-3.076	-3.586	-0.468
MLP (text+vol)	-2.092	-7.189	-18.384	-26.901	MLP (text+vol)	-1.186	-1.571	-2.112	-1.295

Table 24: fincall — temporal split — level-v — R^2_{OOS} vs persistence (test) (n=481).

Table 26: fincall — temporal split — Δv — R^2_{OOS} vs persistence (test) (n=481).

Model	$\tau=3$	$\tau=7$	$\tau=15$	$\tau=30$	Model	$\tau=3$	$\tau=7$	$\tau=15$	$\tau=30$
Ridge (text)	0.416	0.194*	-0.047*	0.112	Ridge (text)	-0.044*	-0.026*	0.008	0.171
MLP (text)	-0.362*	-1.801*	-3.485*	-2.872*	MLP (text)	-0.195*	-0.271*	-0.339*	-0.204*
Ridge (past-vol)	0.432	0.306	0.191*	0.224*	Ridge (past-vol)	0.451	0.306	0.191*	0.224*
MLP (past-vol)	0.482*	0.365*	0.243*	0.212*	MLP (past-vol)	0.492*	0.363*	0.230*	0.249*
Ridge (text+vol)	0.418	0.211	-0.006	0.139	Ridge (text+vol)	0.223*	0.129*	0.094	0.236
MLP (text+vol)	-0.254*	-1.556*	-3.005*	-2.486*	MLP (text+vol)	0.275*	0.140*	-0.061*	0.051

Table 25: fincall — ticker_disjoint split — level-v — R^2_{OOS} vs persistence (test) (n=484).

Table 27: fincall — ticker_disjoint split — Δv — R^2_{OOS} vs persistence (test) (n=484).

A.2 Stage-2 text heads (Result Table 2)

Model	$\tau=3$	$\tau=7$	$\tau=15$	$\tau=30$
Ridge (text)	0.478*	0.225*	-0.033*	-0.322*
MLP (text)	-0.748*	-1.753*	-3.272*	-6.016*
Ridge (past-vol)	0.437	0.222*	0.168	0.208
MLP (past-vol)	0.472*	0.375	0.301*	0.247*
Ridge (text+vol)	0.482*	0.279	0.132	-0.013*
MLP (text+vol)	-0.405*	-1.460*	-2.977*	-5.058*

Table 28: maec — temporal split — level-v — R^2_{OOS} vs persistence (test) (n=507).

Model	$\tau=3$	$\tau=7$	$\tau=15$	$\tau=30$
Ridge (text)	0.466*	0.277	0.117*	-0.036*
MLP (text)	-0.115*	-0.956*	-1.583*	-2.635*
Ridge (past-vol)	0.455*	0.329	0.261	0.262
MLP (past-vol)	0.482*	0.369*	0.281	0.250
Ridge (text+vol)	0.451	0.299	0.186	0.090*
MLP (text+vol)	-0.077*	-0.884*	-1.418*	-2.992*

Table 29: maec — ticker_disjoint split — level-v — R^2_{OOS} vs persistence (test) (n=514).

Model	$\tau=3$	$\tau=7$	$\tau=15$	$\tau=30$
Ridge (text)	0.017*	0.013*	0.047*	-0.074*
MLP (text)	-0.301*	-0.164*	-0.338*	-0.511*
Ridge (past-vol)	0.455	0.313	0.244	0.212
MLP (past-vol)	0.497*	0.395*	0.323*	0.261*
Ridge (text+vol)	-0.005*	-0.142*	0.211	0.138
MLP (text+vol)	0.259*	0.242*	0.068*	-0.080*

Table 30: maec — temporal split — Δv — R^2_{OOS} vs persistence (test) (n=507).

Model	$\tau=3$	$\tau=7$	$\tau=15$	$\tau=30$
Ridge (text)	0.049*	0.056*	0.048*	0.047*
MLP (text)	-0.051*	-0.108*	-0.165*	-0.199*
Ridge (past-vol)	0.455*	0.329	0.261	0.262
MLP (past-vol)	0.489*	0.389*	0.328*	0.294*
Ridge (text+vol)	0.330	0.142*	0.204	0.163
MLP (text+vol)	0.366	0.244	0.120*	-0.111*

Table 31: maec — ticker_disjoint split — Δv — R^2_{OOS} vs persistence (test) (n=514).

Model	$\tau=3$	$\tau=7$	$\tau=15$	$\tau=30$
Ridge (text)	0.659*	0.285	0.200	0.176
MLP (text)	1.001*	0.519*	0.437*	0.405*
Ridge (past-vol)	0.687*	0.859	0.629	0.153*
MLP (past-vol)	0.639*	0.514	0.771	0.387
Ridge (text+vol)	0.648*	0.281	0.199	0.233
MLP (text+vol)	3.631	3.428	4.032	3.109

Table 32: fincall — temporal split — level-v — MSE (test) (n=481).

Model	$\tau=3$	$\tau=7$	$\tau=15$	$\tau=30$
Ridge (text)	0.667	0.311*	0.224*	0.207
MLP (text)	1.556*	1.080*	0.958*	0.903*
Ridge (past-vol)	0.649	0.268	0.173*	0.181*
MLP (past-vol)	0.592*	0.245*	0.162*	0.184*
Ridge (text+vol)	0.665	0.304	0.215	0.201
MLP (text+vol)	1.432*	0.986*	0.856*	0.813*

Table 33: fincall — ticker_disjoint split — level-v — MSE (test) (n=484).

Model	$\tau=3$	$\tau=7$	$\tau=15$	$\tau=30$
Ridge (text)	1.141*	0.406*	0.205	0.127*
MLP (text)	1.302*	0.492*	0.263*	0.228
Ridge (past-vol)	0.687*	0.860	0.803	0.186
MLP (past-vol)	0.717	0.844	0.676	0.196
Ridge (text+vol)	3.145	1.706	0.954	0.164
MLP (text+vol)	2.567	1.076	0.647	0.256

Table 34: fincall — temporal split — Δv — MSE (test) (n=481).

Model	$\tau=3$	$\tau=7$	$\tau=15$	$\tau=30$
Ridge (text)	1.193*	0.396*	0.212	0.193
MLP (text)	1.365*	0.490*	0.286*	0.281*
Ridge (past-vol)	0.627	0.268	0.173*	0.181*
MLP (past-vol)	0.580*	0.246*	0.164*	0.175*
Ridge (text+vol)	0.888*	0.336*	0.194	0.178
MLP (text+vol)	0.828*	0.332*	0.227*	0.221

Table 35: fincall — ticker_disjoint split — Δv — MSE (test) (n=484).

Model	$\tau=3$	$\tau=7$	$\tau=15$	$\tau=30$
Ridge (text)	0.722*	0.443*	0.329*	0.250*
MLP (text)	2.416*	1.573*	1.359*	1.325*
Ridge (past-vol)	0.778	0.445*	0.265	0.150
MLP (past-vol)	0.730*	0.357	0.222*	0.142*
Ridge (text+vol)	0.716*	0.412	0.276	0.191*
MLP (text+vol)	1.943*	1.405*	1.266*	1.144*

Table 36: maec — temporal split — level-v — MSE (test) (n=507).

Model	$\tau=3$	$\tau=7$	$\tau=15$	$\tau=30$
Ridge (text)	0.677*	0.356	0.266*	0.190*
MLP (text)	1.413*	0.963*	0.779*	0.668*
Ridge (past-vol)	0.691*	0.330	0.223	0.136
MLP (past-vol)	0.656*	0.311*	0.217	0.138
Ridge (text+vol)	0.695	0.345	0.246	0.167*
MLP (text+vol)	1.365*	0.927*	0.729*	0.734*

Table 37: maec — ticker_disjoint split — level-v — MSE (test) (n=514).

Model	$\tau=3$	$\tau=7$	$\tau=15$	$\tau=30$
Ridge (text)	1.358*	0.564*	0.303*	0.203*
MLP (text)	1.798*	0.665*	0.426*	0.285*
Ridge (past-vol)	0.753	0.392	0.241	0.149
MLP (past-vol)	0.695*	0.346*	0.215*	0.140*
Ridge (text+vol)	1.390*	0.653*	0.251	0.163
MLP (text+vol)	1.025*	0.433*	0.297*	0.204*

Table 38: maec — temporal split — Δv — MSE (test) (n=507).

Model	$\tau=3$	$\tau=7$	$\tau=15$	$\tau=30$
Ridge (text)	1.206*	0.464*	0.287*	0.175*
MLP (text)	1.332*	0.546*	0.351*	0.220*
Ridge (past-vol)	0.691*	0.330	0.223	0.136
MLP (past-vol)	0.648*	0.301*	0.203*	0.130*
Ridge (text+vol)	0.850	0.422*	0.240	0.154
MLP (text+vol)	0.804	0.372	0.265*	0.204*

Table 39: maec — ticker_disjoint split — Δv — MSE (test) (n=514).

Model	$\tau=3$	$\tau=7$	$\tau=15$	$\tau=30$
ridge eGeMAPS	0.362	0.126*	-0.338*	-1.109*
mlp eGeMAPS	-9.446	-22.924	-39.453	-63.845
ridge eGeMAPS +vol	0.347	0.258	-0.046	-0.668
mlp eGeMAPS +vol	-9.226	-21.289	-46.203	-77.488
ridge WavLM	0.359	0.122*	-0.427*	-1.167*
mlp WavLM	-0.117*	-0.965*	-2.387*	-5.448*
ridge WavLM +vol	0.343	0.245	-0.390	-1.599
mlp WavLM +vol	-4.487	-14.363	-32.806	-59.317
ridge emotion2vec+	0.353	0.138*	-0.289*	-1.025*
mlp emotion2vec+	-0.385*	-1.880*	-4.314*	-7.784*
ridge emotion2vec+ +vol	0.401	-0.583	-1.139	-0.422*
mlp emotion2vec+ +vol	-6.056	-20.286	-41.247	-64.876
ridge WavLM+eGeMAPS	0.358	0.116*	-0.379*	-1.172*
mlp WavLM+eGeMAPS	-2.476	-6.938	-11.288	-17.515
ridge WavLM+eGeMAPS +vol	0.338	0.221	-0.139	-1.661
mlp WavLM+eGeMAPS +vol	-5.037	-14.919	-30.138	-43.263
ridge WavLM+text +vol	0.443*	0.321	0.014	-1.142
mlp WavLM+text +vol	-2.902	-8.912	-22.030	-33.025

Table 40: fincall — temporal split — level-v — R^2_{OOS} vs persistence (test) (n=481).

Model	$\tau=3$	$\tau=7$	$\tau=15$	$\tau=30$
ridge eGeMAPS	0.410	0.106*	-0.238*	-0.106*
mlp eGeMAPS	-0.074	-1.156*	-1.969*	-1.835*
ridge eGeMAPS +vol	0.416	0.270	0.164*	0.106*
mlp eGeMAPS +vol	-0.338	-1.497	-2.709	-2.606*
ridge WavLM	0.397	0.092*	-0.235*	-0.092*
mlp WavLM	-0.355*	-1.902*	-4.384*	-4.021*
ridge WavLM +vol	0.403	0.137*	0.088	0.106*
mlp WavLM +vol	-0.320*	-1.388*	-2.985*	-3.161*
ridge emotion2vec+	0.409	0.125*	-0.202*	-0.061*
mlp emotion2vec+	-1.476	-5.431	-10.233	-9.811*
ridge emotion2vec+ +vol	0.416	0.298	0.195*	0.211*
mlp emotion2vec+ +vol	-1.243	-4.109	-8.681	-7.141*
ridge WavLM+eGeMAPS	0.394	0.094*	-0.230*	-0.092*
mlp WavLM+eGeMAPS	-0.308*	-1.393*	-2.760*	-2.331*
ridge WavLM+eGeMAPS +vol	0.400	0.137*	0.073	0.106*
mlp WavLM+eGeMAPS +vol	-0.302*	-1.323*	-2.551*	-2.151*
ridge WavLM+text +vol	0.412	0.205	0.001	0.106*
mlp WavLM+text +vol	-0.136*	-1.278*	-2.540*	-2.041*

Table 41: fincall — ticker_disjoint split — level-v — R^2_{OOS} vs persistence (test) (n=484).

A.3 Stage-3 audio heads (Result Table 3)

FinCall-Surprise only (MAEC ships no raw audio).

Model	$\tau=3$	$\tau=7$	$\tau=15$	$\tau=30$
ridge eGeMAPS	-0.000*	-0.007*	-0.015	-0.000*
mlp eGeMAPS	-0.492*	-0.218*	-0.272*	-0.335*
ridge eGeMAPS +vol	0.219	-1.834	-4.099	-0.611*
mlp eGeMAPS +vol	0.187	-1.680	-5.488	-5.311*
ridge WavLM	-0.005*	-0.007*	-0.045	-0.031*
mlp WavLM	-0.176*	-0.248*	-0.449*	-0.911*
ridge WavLM +vol	0.265	-1.026	-2.032	-0.511*
mlp WavLM +vol	-1.219	-2.099	-2.908	-2.411*
ridge emotion2vec+	-0.018*	-0.028*	-0.056	-0.000*
mlp emotion2vec+	-0.096*	-0.256*	-0.468*	-1.241*
ridge emotion2vec+ +vol	0.229	-2.052	-3.031	-0.611*
mlp emotion2vec+ +vol	-1.692	-3.466	-3.611	-2.311*
ridge WavLM+eGeMAPS	-0.005*	0.001*	-0.015	-0.021*
mlp WavLM+eGeMAPS	-0.296*	-0.497*	-0.641*	-1.341*
ridge WavLM+eGeMAPS +vol	0.214	-1.099	-3.706	-0.511*
mlp WavLM+eGeMAPS +vol	-1.810	-3.234	-3.991	-3.111*
ridge WavLM+text +vol	-1.697	-3.043	-0.323	-0.411*
mlp WavLM+text +vol	-1.456	-2.193	-3.917	-3.111*

Table 42: fincall — temporal split — Δv — R^2_{OOS} vs persistence (test) (n=481).

Model	$\tau=3$	$\tau=7$	Model	$\tau=30$	$\tau=3$	$\tau=7$	$\tau=15$	$\tau=30$
ridge eGeMAPS	-0.011*	-0.013*	ridge eGeMAPS	0.002	0.674	0.345*	0.265*	0.258*
mlp eGeMAPS	-0.167*	-0.231*	mlp eGeMAPS	-0.171*	1.227	0.832*	0.634*	0.666*
ridge eGeMAPS +vol	0.415	0.253	ridge eGeMAPS +vol	0.002	0.667	0.281	0.179*	0.188
mlp eGeMAPS +vol	0.353	0.207	mlp eGeMAPS +vol	0.070	1.529	0.963	0.792	0.843
ridge WavLM	-0.038*	-0.035*	ridge WavLM	-0.027*	0.689	0.350*	0.264*	0.256*
mlp WavLM	-0.248*	-0.323*	mlp WavLM	-0.351*	1.547*	1.119*	1.150*	1.172*
ridge WavLM +vol	0.337*	0.171*	ridge WavLM +vol	0.024	0.682	0.333*	0.195	0.198
mlp WavLM +vol	0.160*	0.009*	mlp WavLM +vol	-0.174*	1.508*	0.921*	0.852*	0.971*
ridge emotion2vec+	-0.002*	-0.003*	ridge emotion2vec+	0.002	0.675	0.338*	0.257*	0.249*
mlp emotion2vec+	-0.257*	-0.397*	mlp emotion2vec+	-0.492*	2.828	2.480	2.400	2.539
ridge emotion2vec+ +vol	0.435	0.297	ridge emotion2vec+ +vol	0.108	0.667	0.271	0.172*	0.182*
mlp emotion2vec+ +vol	0.105*	-0.139*	mlp emotion2vec+ +vol	-0.255*	2.562	1.970	2.068	1.909
ridge WavLM+eGeMAPS	-0.050*	-0.042*	ridge WavLM+eGeMAPS	-0.026*	0.692	0.350*	0.263*	0.256*
mlp WavLM+eGeMAPS	-0.225*	-0.366*	mlp WavLM+eGeMAPS	-0.451*	1.493*	0.923*	0.803*	0.778*
ridge WavLM+eGeMAPS +vol	0.304*	0.147*	ridge WavLM+eGeMAPS +vol	0.020	0.686	0.333*	0.198	0.200
mlp WavLM+eGeMAPS +vol	0.109*	-0.004*	mlp WavLM+eGeMAPS +vol	-0.111*	1.487*	0.896*	0.759*	0.736*
ridge WavLM+text +vol	0.159*	0.072*	ridge WavLM+text +vol	0.023	0.671	0.307	0.214	0.199
mlp WavLM+text +vol	0.228*	0.096*	mlp WavLM+text +vol	-0.084*	1.298*	0.878*	0.756*	0.709*

Table 43: fincall — ticker_disjoint split — Δv — R^2_{OOS} vs persistence (test) (n=484).

Table 45: fincall — ticker_disjoint split — level-v — MSE (test) (n=484).

Model	$\tau=3$	$\tau=7$	Model	$\tau=30$	$\tau=3$	$\tau=7$	$\tau=15$	$\tau=30$
ridge eGeMAPS	0.750	0.366*	ridge eGeMAPS	0.278*	1.175*	0.421*	0.211	0.112*
mlp eGeMAPS	12.267	10.015	mlp eGeMAPS	8.416	1.752*	0.510*	0.265*	0.149*
ridge eGeMAPS +vol	0.767	0.311	ridge eGeMAPS +vol	0.218	0.917	1.186	1.061	0.182
mlp eGeMAPS +vol	12.008	9.331	mlp eGeMAPS +vol	9.819	0.954	1.122	1.350	0.704
ridge WavLM	0.753	0.368*	ridge WavLM	0.297*	1.180*	0.422*	0.217	0.115*
mlp WavLM	1.312*	0.823*	mlp WavLM	0.704	1.381*	0.523*	0.301*	0.221
ridge WavLM +vol	0.772	0.316	ridge WavLM +vol	0.289	0.863	0.848	0.631	0.174
mlp WavLM +vol	6.443	6.432	mlp WavLM +vol	7.032	2.606	1.297	0.813	0.389
ridge emotion2vec+	0.759	0.361*	ridge emotion2vec+	0.226	1.196*	0.431*	0.220	0.118*
mlp emotion2vec+	1.626*	1.206*	mlp emotion2vec+	0.979	1.288*	0.526*	0.305*	0.250*
ridge emotion2vec+ +vol	0.703	0.663	ridge emotion2vec+ +vol	0.158	0.905	1.278	0.838	0.187
mlp emotion2vec+ +vol	8.286	8.911	mlp emotion2vec+ +vol	7.340	3.161	1.870	0.959	0.370
ridge WavLM+eGeMAPS	0.754	0.370*	ridge WavLM+eGeMAPS	0.287*	1.180*	0.418*	0.211	0.115*
mlp WavLM+eGeMAPS	4.082	3.323	mlp WavLM+eGeMAPS	2.556	1.522*	0.627*	0.341*	0.261*
ridge WavLM+eGeMAPS +vol	0.777	0.326	ridge WavLM+eGeMAPS +vol	0.237	0.923	0.879	0.979	0.172
mlp WavLM+eGeMAPS +vol	7.089	6.664	mlp WavLM+eGeMAPS +vol	6.477	3.300	1.773	1.038	0.466
ridge WavLM+text +vol	0.654*	0.284	ridge WavLM+text +vol	0.205	3.167	1.692	0.275	0.161
mlp WavLM+text +vol	4.582	4.150	mlp WavLM+text +vol	4.791	2.884	1.337	1.023	0.458

Table 44: fincall — temporal split — level-v — MSE (test) (n=481).

Table 46: fincall — temporal split — Δv — MSE (test) (n=481).

Model	$\tau=3$	$\tau=7$	$\tau=15$	$\tau=30$
ridge eGeMAPS	1.155*	0.391*	0.214	0.233
mlp eGeMAPS	1.333*	0.475*	0.251*	0.258*
ridge eGeMAPS +vol	0.668	0.288	0.180	0.186*
mlp eGeMAPS +vol	0.739	0.306	0.199	0.217
ridge WavLM	1.185*	0.399*	0.219*	0.228
mlp WavLM	1.426*	0.510*	0.289*	0.275*
ridge WavLM +vol	0.758*	0.320*	0.190	0.184*
mlp WavLM +vol	0.959*	0.382*	0.251*	0.261*
ridge emotion2vec+	1.144*	0.387*	0.211	0.233
mlp emotion2vec+	1.436*	0.539*	0.319*	0.335*
ridge emotion2vec+ +vol	0.645	0.271	0.171*	0.182*
mlp emotion2vec+ +vol	1.022*	0.439*	0.268*	0.275*
ridge WavLM+eGeMAPS	1.199*	0.402*	0.219*	0.227
mlp WavLM+eGeMAPS	1.400*	0.527*	0.312*	0.316*
ridge WavLM+eGeMAPS +vol	0.795*	0.329*	0.204	0.184*
mlp WavLM+eGeMAPS +vol	1.017*	0.387*	0.239*	0.260*
ridge WavLM+text +vol	0.961*	0.358*	0.207	0.186
mlp WavLM+text +vol	0.882*	0.348*	0.232*	0.216

Table 47: fincall — ticker_disjoint split — Δv — MSE (test) (n=484).

Stage	$\tau=3$	$\tau=7$	$\tau=15$	$\tau=30$	Stage	$\tau=3$	$\tau=7$	$\tau=15$	$\tau=30$
Persistence	0.000	0.000	0.000	0.000	Persistence	0.000	0.000	0.000	0.000
HAR-RV	0.464	0.410	0.207	-0.287	HAR-RV	0.464	0.410	0.207	-0.287
GBDT (ticker FE)	0.354	0.262	-0.040	-0.839	GBDT (ticker FE)	0.344	0.248	-0.057	-0.904
Text (BGE+vol)	0.448	0.328	0.045	-1.091	Text (BGE+vol)	-1.678	-3.076	-3.586	-0.468
Audio (WavLM+eG+vol)	0.338	0.221	-0.139	-1.661	Audio (WavLM+eG+vol)	0.214	-1.099	-3.706	-0.548
Fusion (stack)	0.448	0.370	0.125	-1.045	Fusion (stack)	0.232	-0.326	-1.828	-0.010*

Table 48: fincall — temporal split — level-v — R^2_{OOS} vs persistence (test) (n=481).

Table 50: fincall — temporal split — Δv — R^2_{OOS} vs persistence (test) (n=481).

Stage	$\tau=3$	$\tau=7$	$\tau=15$	$\tau=30$	Stage	$\tau=3$	$\tau=7$	$\tau=15$	$\tau=30$
Persistence	0.000	0.000	0.000	0.000	Persistence	0.000	0.000	0.000	0.000
HAR-RV	0.495	0.357	0.211	0.229	HAR-RV	0.495	0.357	0.211	0.229
GBDT (ticker FE)	0.412	0.278	0.078	0.116	GBDT (ticker FE)	0.407	0.275	0.089	0.108
Text (BGE+vol)	0.418	0.211	-0.006	0.133	Text (BGE+vol)	0.223*	0.129*	0.094	0.236
Audio (WavLM+eG+vol)	0.400	0.137	0.073	0.144	Audio (WavLM+eG+vol)	0.304*	0.147*	0.044	0.209
Fusion (stack)	0.450	0.264	0.050	0.171	Fusion (stack)	0.436	0.292	0.133	0.244*

Table 49: fincall — ticker_disjoint split — level-v — R^2_{OOS} vs persistence (test) (n=484).

Table 51: fincall — ticker_disjoint split — Δv — R^2_{OOS} vs persistence (test) (n=484).

A.4 Consolidated ablation grid (Result Table 4)

One pre-registered canonical model per stage; * marks Holm-corrected significance vs. Stage-1 within the Δv confirmatory family (two-sided—read direction from the value).

Stage	$\tau=3$	$\tau=7$	$\tau=15$	$\tau=30$
Persistence	0.000	0.000	0.000	0.000
HAR-RV	0.507	0.392	0.306	0.206
GBDT (ticker FE)	0.397	0.333	0.217	0.131
Text (BGE+vol)	0.482	0.279	0.132	-0.013
Audio (WavLM+eG+vol)	—	—	—	—
Fusion (stack)	—	—	—	—

Table 52: maec — temporal split — level-v — R^2_{OOS} vs persistence (test) (n=507).

Stage	$\tau=3$	$\tau=7$	$\tau=15$	$\tau=30$
Persistence	0.000	0.000	0.000	0.000
HAR-RV	0.474	0.379	0.311	0.225
GBDT (ticker FE)	0.395	0.306	0.236	0.221
Text (BGE+vol)	0.451	0.299	0.186	0.090
Audio (WavLM+eG+vol)	—	—	—	—
Fusion (stack)	—	—	—	—

Table 53: maec — ticker_disjoint split — level-v — R^2_{OOS} vs persistence (test) (n=514).

Stage	$\tau=3$	$\tau=7$	$\tau=15$	$\tau=30$
Persistence	0.000	0.000	0.000	0.000
HAR-RV	0.507	0.392	0.306	0.206
GBDT (ticker FE)	0.407	0.340	0.213	0.186
Text (BGE+vol)	-0.005*	-0.142*	0.211	0.138
Audio (WavLM+eG+vol)	—	—	—	—
Fusion (stack)	—	—	—	—

Table 54: maec — temporal split — Δv — R^2_{OOS} vs persistence (test) (n=507).

Stage	$\tau=3$	$\tau=7$	$\tau=15$	$\tau=30$
Persistence	0.000	0.000	0.000	0.000
HAR-RV	0.474	0.379	0.311	0.225
GBDT (ticker FE)	0.394	0.285	0.227	0.233
Text (BGE+vol)	0.330	0.142*	0.204	0.163
Audio (WavLM+eG+vol)	—	—	—	—
Fusion (stack)	—	—	—	—

Table 55: maec — ticker_disjoint split — Δv — R^2_{OOS} vs persistence (test) (n=514).

Stage	$\tau=3$	$\tau=7$	$\tau=15$	$\tau=30$
Persistence	1.383	0.571	0.318	0.189
HAR-RV	0.682	0.347	0.221	0.150
GBDT (ticker FE)	0.834	0.381	0.249	0.164
Text (BGE+vol)	0.716	0.412	0.276	0.191
Audio (WavLM+eG+vol)	—	—	—	—
Fusion (stack)	—	—	—	—

Table 60: maec — temporal split — level-v — MSE (test) (n=507).

Stage	$\tau=3$	$\tau=7$	$\tau=15$	$\tau=30$
Persistence	1.174	0.419	0.208	0.111
HAR-RV	0.630	0.247	0.165	0.143
GBDT (ticker FE)	0.758	0.309	0.216	0.205
Text (BGE+vol)	0.648	0.281	0.199	0.233
Audio (WavLM+eG+vol)	0.777	0.326	0.237	0.296
Fusion (stack)	0.649	0.264	0.182	0.225

Table 56: fincall — temporal split — level-v — MSE (test) (n=481).

Stage	$\tau=3$	$\tau=7$	$\tau=15$	$\tau=30$
Persistence	1.268	0.492	0.302	0.184
HAR-RV	0.667	0.306	0.208	0.142
GBDT (ticker FE)	0.767	0.342	0.230	0.143
Text (BGE+vol)	0.695	0.345	0.246	0.167
Audio (WavLM+eG+vol)	—	—	—	—
Fusion (stack)	—	—	—	—

Stage	$\tau=3$	$\tau=7$	$\tau=15$	$\tau=30$
Persistence	1.142	0.386	0.214	0.233
HAR-RV	0.577	0.248	0.169	0.180
GBDT (ticker FE)	0.672	0.278	0.197	0.206
Text (BGE+vol)	0.665	0.304	0.215	0.201
Audio (WavLM+eG+vol)	0.686	0.333	0.198	0.200
Fusion (stack)	0.628	0.284	0.203	0.193

Table 57: fincall — ticker_disjoint split — level-v — MSE (test) (n=484).

Table 61: maec — ticker_disjoint split — level-v — MSE (test) (n=514).

Stage	$\tau=3$	$\tau=7$	$\tau=15$	$\tau=30$
Persistence	1.174	0.419	0.208	0.111
HAR-RV	0.630	0.247	0.165	0.143
GBDT (ticker FE)	0.771	0.315	0.220	0.212
Text (BGE+vol)	3.145	1.706	0.954	0.164
Audio (WavLM+eG+vol)	0.923	0.879	0.979	0.172
Fusion (stack)	0.902	0.555	0.588	0.113*

Stage	$\tau=3$	$\tau=7$	$\tau=15$	$\tau=30$
Persistence	1.383	0.571	0.318	0.189
HAR-RV	0.682	0.347	0.221	0.150
GBDT (ticker FE)	0.820	0.377	0.250	0.154
Text (BGE+vol)	1.390*	0.653*	0.251	0.163
Audio (WavLM+eG+vol)	—	—	—	—
Fusion (stack)	—	—	—	—

Table 62: maec — temporal split — Δv — MSE (test) (n=507).

Table 58: fincall — temporal split — Δv — MSE (test) (n=481).

Stage	$\tau=3$	$\tau=7$	$\tau=15$	$\tau=30$
Persistence	1.142	0.386	0.214	0.233
HAR-RV	0.577	0.248	0.169	0.180
GBDT (ticker FE)	0.678	0.279	0.195	0.208
Text (BGE+vol)	0.888*	0.336*	0.194	0.178
Audio (WavLM+eG+vol)	0.795*	0.329*	0.204	0.184
Fusion (stack)	0.644	0.273	0.185	0.176*

Stage	$\tau=3$	$\tau=7$	$\tau=15$	$\tau=30$
Persistence	1.268	0.492	0.302	0.184
HAR-RV	0.667	0.306	0.208	0.142
GBDT (ticker FE)	0.769	0.352	0.233	0.141
Text (BGE+vol)	0.850	0.422*	0.240	0.154
Audio (WavLM+eG+vol)	—	—	—	—
Fusion (stack)	—	—	—	—

Table 63: maec — ticker_disjoint split — Δv — MSE (test) (n=514).

Table 59: fincall — ticker_disjoint split — Δv — MSE (test) (n=484).

A.5 Substrate audit (Result Table 5R)

[pending: T6R.1: the full per-file audit behind Table 1—SHA-256 manifests of the fetched Volt-AGE/KeFVP/HTML label and split files, boundary dates and call counts, overlap by cohort, the sentinel-zero cells, the interleaved-column schema hazard, and the diff of the shipped labels against our recomputed targets (MAEC join verified at 99.4%).]

A.6 Reproductions and code-availability audit (Result Table 6R)

[pending: T6R.2: the complete reproduction grid behind Table 7, and the audit of the fifteen non-runnable papers—repository URL, HTTP status and date, and the reason code (no code; empty or unpushed repository; closed-API dependency; stated inability to run; paywalled).]